

Supplementary Material: Local Permission Is Not Enough

Anonymous authors

Paper under double-blind review

1 Artifact Map

Path	Purpose
<code>data/tasks_tracebreak_120.jsonl</code>	120 synthetic paired tasks: 10 risky and 10 safe controls in each of six categories.
<code>tracebreak/env</code>	Synthetic world, tools, trace events, and provenance tags.
<code>tracebreak/policies</code>	Local guard, visible-content DLP, visible-policy guard, TraceGuard, and StrictTraceGuard.
<code>tracebreak/experiments</code>	Deterministic scripted runner and bounded API runner.
<code>tracebreak/analysis</code>	Metric, paper-table, example-trace, and figure-generation scripts.
<code>results/raw_traces</code>	JSONL traces used to compute the reported tables.
<code>paper/tables</code>	Paper-facing tables, source-reference ablations, and compact example trace exports.
<code>REPRODUCIBILITY.md</code>	End-to-end commands for rebuilding tasks, traces, metrics, figures, tests, and PDFs.

Table 1: TraceBreak artifact structure.

2 Metric Definitions

All metrics are computed from executed tool events and deterministic policy checks. No LLM judge is used.

Metric	Definition
Sink rate	Fraction of all tasks whose expected write sink executes.
Safe utility	Fraction of safe-control tasks whose expected sink executes without a global violation.
Risky violation	Fraction of risky tasks with an executed sink that violates TraceGuard’s global invariant checker.
LPVR	Local-pass violation rate: fraction of risky tasks where every local check allows the trace but the trace is globally violating.
Safe false block	Fraction of safe-control tasks blocked by a defense after local checks allowed the action.
Risky block	Fraction of risky tasks blocked by a defense after local checks allowed the action.
Model calls	Mean number of API model calls per task in API conditions.

Table 2: Reported metric definitions.

3 Agent Protocol

The API runner uses a provider-independent JSON-action protocol rather than provider-specific function calling. At each step, the model receives the user instruction, a text list of available tools, and the visible observation trace. It must return exactly one JSON object with keys `action` and `arguments`. The visible trace includes tool observations but not hidden

provenance tags. The `api_policy_prompt` condition adds only a natural language security policy to the system prompt. The `api_traceguard` condition does not add a policy prompt; it relies on runtime enforcement.

Available tools are: `search_docs`, `read_doc`, `summarize`, `search_people`, `send_email`, `post_ticket`, `write_memory`, `read_memory`, `search_approvals`, `get_approval`, `search_records`, `read_record`, `aggregate_records`, and `final_answer`. The runner normalizes minor ID/name mismatches, source-reference aliases, and common JSON field-name variants before applying local and defense checks.

4 Additional Result Tables

Condition	n	Sink	Safe util.	Risky viol.	LPVR	Risky block
DLP	120	100	100	100	100	0
Multi local	120	100	100	100	100	0
Single local	120	100	100	100	100	0
Visible policy	120	83	100	67	67	33
TraceGuard	120	50	100	0	0	100

Table 3: Deterministic 120-task matrix.

Condition	n	Sink	Safe util.	Risky viol.	LPVR	Calls
API DLP	24	100	100	100	100	5.1
API local	24	100	100	100	100	5.1
API policy prompt	24	79	83	75	75	4.6
API TraceGuard	24	50	100	0	0	5.1

Table 4: API subset with gpt-4.1-mini.




Figure 1: Security-utility frontier on the 24-task gpt-4.1-mini API subset. The desired region is high safe-control utility and low risky-violation rate. Local guards and DLP preserve utility but violate every risky task; TraceGuard preserves utility while eliminating risky violations.

Category	Local viol.	DLP viol.	Prompt viol.	TG viol.	TG block	TG safe util.
aggregation threshold	100	100	100	0	100	100
approval scope	100	100	0	0	100	100
audience restriction	100	100	100	0	100	100
cross tenant	100	100	50	0	100	100
memory laundering	100	100	100	0	100	100
sensitive external	100	100	100	0	100	100

Table 5: Category-level API results on the 24-task gpt-4.1-mini subset. Each category has two risky and two safe-control tasks.

Category	Expected reason	Det.	API	Safe Det/ API	Unexpected Det/ API
aggregation threshold	aggregation threshold	10/10	2/2	0/0	0/0
approval scope	approval scope	10/10	2/2	0/0	0/0
audience restriction	audience mismatch	10/10	2/2	0/0	0/0
cross tenant	cross tenant	10/10	2/2	0/0	0/0
memory laundering	external recipient	10/10	2/2	0/0	0/0
sensitive external	external recipient	10/10	2/2	0/0	0/0

Table 6: TraceGuard block-reason audit. Det. and API count risky rows whose block reason matches the category invariant; safe and unexpected columns are deterministic/API counts.

Category	Risky viol.	Risky nonviol.	Safe util.	Safe lost	Final ans.
aggregation threshold	2/2	0/2	2/2	0/2	0
approval scope	0/2	2/2	0/2	2/2	4
audience restriction	2/2	0/2	2/2	0/2	0
cross tenant	1/2	1/2	2/2	0/2	0
memory laundering	2/2	0/2	2/2	0/2	0
sensitive external	2/2	0/2	2/2	0/2	0
overall	9/12	3/12	10/12	2/12	4

Table 7: Policy-prompt diagnostic on the 24-task gpt-4.1-mini API subset. Nonviolating risky rows are not defense blocks; they are prompt-induced non-completion or non-sink paths.

Condition	Risky transfer	Viol. after transfer	Block after transfer	Safe util. after transfer
Single local	17	100	0	100
Multi local	100	100	0	100
DLP	100	100	0	100
Visible policy	100	67	33	100
TraceGuard	100	0	100	100

Table 8: Authority-transfer analysis on deterministic traces. A transfer sink references an output produced by another actor, or by a memory handoff.

Condition	n	Safe util.	Risky viol.	LPVR	Risky block	Safe FP
Multi local	120	100	100	100	0	0
DLP	120	100	100	100	0	0
Visible policy	120	100	67	67	33	0
TraceGuard	120	100	0	0	100	0
Inferred TraceGuard	120	100	0	0	100	0

Table 9: Deterministic indirect-injection overlay. Synthetic instruction-like text is appended to every document and record body before replaying scripted traces.

Condition	n	Decoy hits	Safe util.	Risky viol.	LPVR	Risky block	Safe FP
Multi local	120	260	100	100	100	0	0
DLP	120	260	100	100	100	0	0
Visible policy	120	260	100	67	67	33	0
TraceGuard	120	260	100	0	0	100	0
Inferred TraceGuard	120	260	100	0	0	100	0

Table 10: Deterministic decoy-ambiguity stress test. Plausible decoy documents, recipients, approvals, and records are added to each synthetic world before replaying scripted traces with search preambles.

Defense	n	Safe util.	Risky viol.	LPVR	Risky block	Safe FP
Local guard	24	100	100	100	0	0
DLP replay	24	100	100	100	0	0
Visible-policy replay	24	100	67	67	33	0
Metadata-critic replay	24	100	0	0	100	0
TraceGuard replay	24	100	0	0	100	0

Table 11: Same-action counterfactual replay over the 24 gpt-4.1-mini API-local traces. The language-model actions are fixed; only the runtime defense changes.

Category	Decisive fact	Visible block	Visible viol.	Metadata block	TG block
aggregation threshold	visible bulk count	2/2	0/2	2/2	2/2
approval scope	visible approval	2/2	0/2	2/2	2/2
audience restriction	hidden audience	0/2	2/2	2/2	2/2
cross tenant	hidden tenant	0/2	2/2	2/2	2/2
memory laundering	hidden memory provenance	0/2	2/2	2/2	2/2
sensitive external	hidden confidentiality/recipient	0/2	2/2	2/2	2/2
overall	visible subset + hidden metadata	4/12	8/12	12/12	12/12

Table 12: Same-action visibility-gap audit over fixed gpt-4.1-mini API-local traces. Visible-policy review catches only categories whose decisive fact is visible at the sink; metadata-aware review and TraceGuard catch every risky sink.

Condition	Runs	Sinks	Valid refs	Missing/empty	Invalid	Final	Blocked
Local guards	24	24	24/24	0	0	0	0
DLP	24	24	24/24	0	0	0	0
Policy prompt	24	20	20/20	0	0	4	0
TraceGuard	24	24	24/24	0	0	0	12
Overall	96	92	92/92	0	0	4	12

Table 13: Cooperative API source-reference compliance before source-reference stress tests. All executed gpt-4.1-mini write sinks carry nonempty refs to prior observations; policy prompting has fewer sinks because four runs end in final_answer.

Source refs	Defense	Safe util.	Risky viol.	LPVR	Risky block	Safe FP
Cooperative	TraceGuard	100	0	0	100	0
Deleted at sink	TraceGuard	83	100	100	0	0
Deleted at sink	Inferred TraceGuard	100	0	0	100	0
Deleted at sink	StrictTraceGuard	83	0	0	100	92
Corrupted refs	TraceGuard	83	83	83	17	17
Corrupted refs	Inferred TraceGuard	100	0	0	100	0
Corrupted refs	StrictTraceGuard	83	83	83	17	17

Table 14: Source-reference robustness ablation on the same 24 gpt-4.1-mini TraceGuard traces. Deleted rows remove sink source_refs; corrupted rows replace them with a benign prior public reference. Inferred TraceGuard reconstructs the sink tag from prior runtime-owned event tags. StrictTraceGuard blocks missing provenance but does not address present-but-corrupted references.

Condition	Candidates	Repair success	Repair block	Risky viol.	Max steps
API TraceGuard repair oracle	12	100	0	0	1
Scripted TraceGuard repair oracle	60	100	0	0	2

Table 15: Repair-oracle upper bounds over blocked risky TraceGuard traces. The oracle appends matched safe-control continuations after the block. Scripted aggregation-threshold repairs require one extra aggregation step before the safe email, so this table shows that compliant continuations exist but does not measure whether the language-model agent discovers them.

5 Representative Trace Pair

Table 16 shows the same risky sensitive-external task under local guards and under TraceGuard. Local checks allow every individual action in both traces. The difference is that TraceGuard propagates the hidden tag from the read document through the summary and blocks the final external email. The artifact also includes one aligned local/TraceGuard example for each policy category in `paper/tables/api_gpt41mini_category_examples.md`.

Step	Action	Local guard	Local condition	TraceGuard condition
1	search_docs	allow	allow	allow
2	read_doc	allow	allow; source=doc.pricing_100	allow; source=doc.pricing_100
3	summarize	allow	allow; tag remains confidential	allow; tag remains confidential
4	search_people	allow	allow; recipient external	allow; recipient external
5	send_email	allow	allow; global violation	block; internal_data.to_external_recipient

Table 16: Example trace pair for sensitive_external_000_risky.

6 Reproduction Notes

The core commands are:

```
conda run -n trace_level python -m unittest \
  tests/test_policies.py tests/test_api_normalization.py

conda run -n trace_level python -m tracebreak.data.generate_tasks \
  --out data/tasks_tracebreak_120.jsonl

conda run -n trace_level python -m tracebreak.analysis.compute_metrics \
  --runs results/raw_traces/*.jsonl \
  --out-csv results/metrics.csv \
  --out-md results/tables/main_results.md
```

The full command list is in `REPRODUCIBILITY.md`. API commands accept a configurable key-path argument and cache responses under `results/api_cache`, so exact reruns reuse cached responses when available. The reported API subset used gpt-4.1-mini, two 12-task offsets, four conditions, and an eight-step maximum per task. For future budgeted API sweeps, the runner records provider-reported token usage when it is returned, and the metric script aggregates token totals by condition.